

# What is Data Retrieval?



**SDSC** SAN DIEGO  
SUPERCOMPUTER CENTER

# What is Data Retrieval?

- **Data retrieval**

- The way in which the desired data is specified and retrieved from a data store

- **Our focus**

- ① • How to specify a data request (query specification method)
  - For static and streaming data < big data ★ (focus on this)

- ② • The internal mechanism of data retrieval

- For large and streaming data

# What is a Query Language?

- A language to specify the data items you need
- A query language is declarative
  - Specify ~~what~~ you need rather than ~~how~~ to obtain it  
(you specify and the system does the rest)
  - SQL (Structured Query Language) ~ most used query language for relational data
    - Database programming language — like Oracle's PL/SQL, PostgreSQL's pgSQL
    - Procedural programming language
      - Embeds query operations

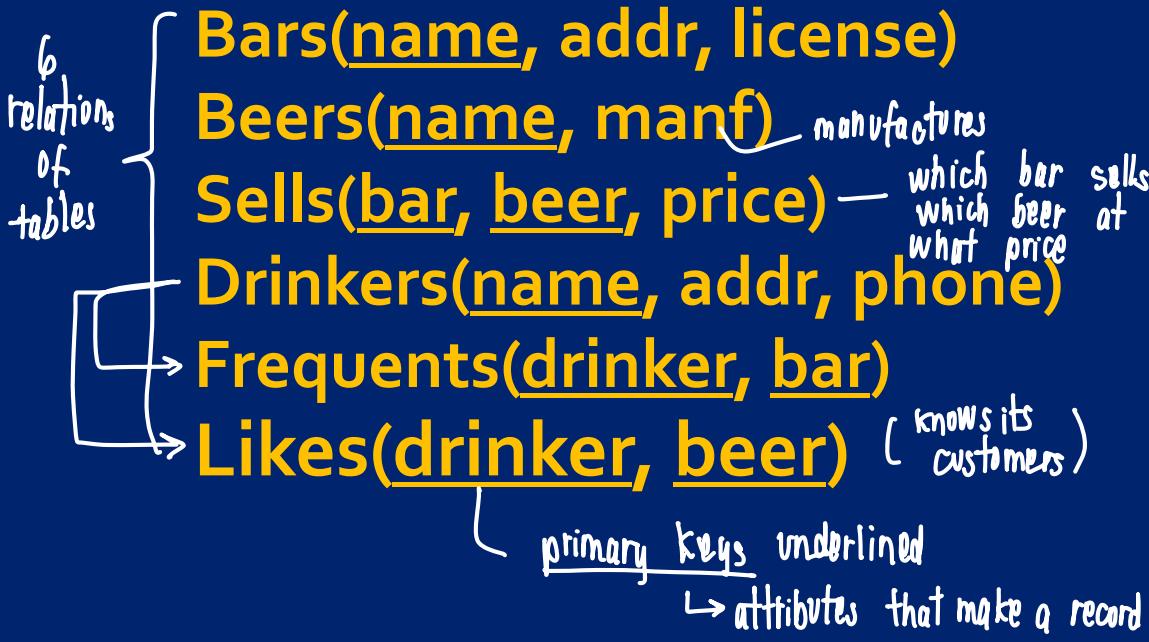
① **SQL** — ubiquitous query language when data is structured ~ has been extended to accommodate other types of data

← oracle  
spark

- The standard for structured data

- Oracle's SQL to Spark SQL

- Example Database Schema



<u>name</u>	<u>addr</u>	license
Great American Bar	363 Main St., SD, CA 92390	41-437844098
Beer Paradise	6450 Mango Drive, SD, CA 92130	41-973428319
Have a Good Time	8236 Adams Avenue, SD, CA 92116	32-032263401

# SELECT-FROM-WHERE

- Which beers are made by Heineken?

*SELECT name  
FROM Beers  
WHERE manf='Heineken'*

\* match exactly  
single string literal

The condition(s) to satisfy

Strings like 'Heineken' are case-sensitive and are put in quotes

Output attribute(s)

Table(s) to use

name
Heineken Lager Beer
Amstel Lager
Amstel Light
...

Select *manf='Heineken'* (Beers)



Project(*name*)

result  
table with  
name row

# More Example Queries

- **Find expensive beer**

- SELECT DISTINCT beer, price
  - FROM Sells
  - WHERE price > 15
- ensures that result relation will have no duplicate*

- **Which businesses have a Temporary License (starts with 32) in San Diego?**

- SELECT name
- FROM Bars
- WHERE addr LIKE '%SD%' **AND** license LIKE '32%' **LIMIT** 5

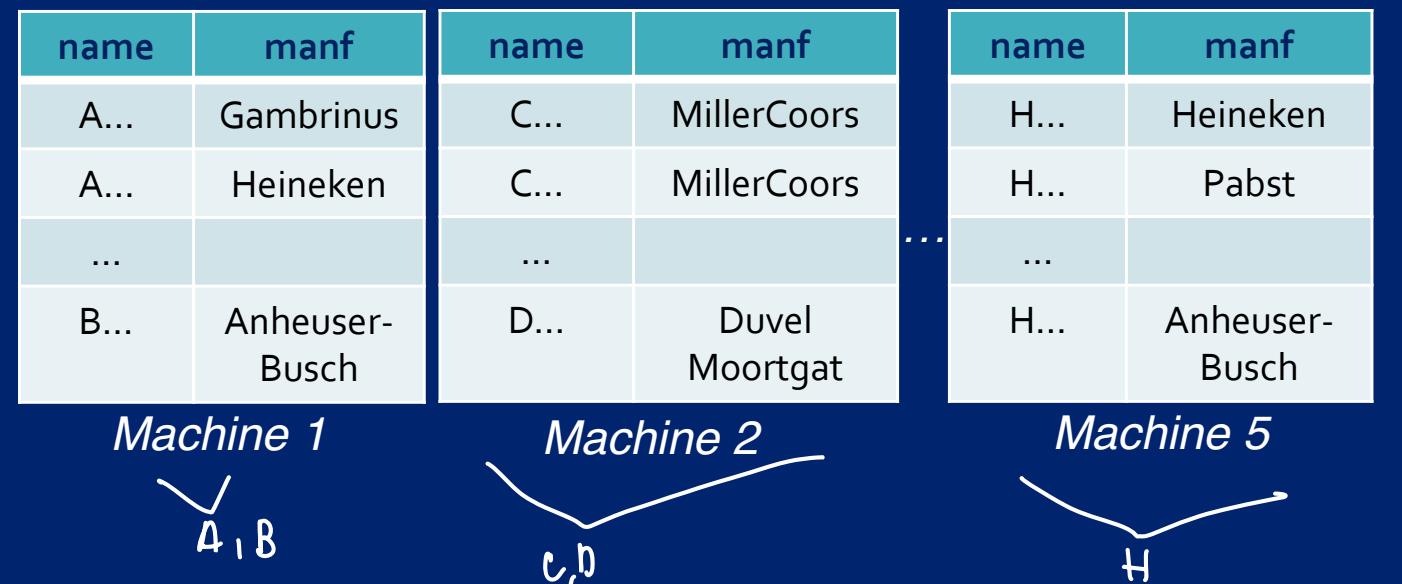
*limit on number of results to return*

*\* limit clause may vary between DBMS vendors*

<u>name</u>	<u>addr</u>	license
Great American Bar	363 Main St., SD, CA 92390	41-437844098
Beer Paradise	6450 Mango Drive, SD, CA 92130	41-973428319
Have a Good Time	8236 Adams Avenue, SD, CA 92116	32-032263401

# Select-Project Queries in the Large

- Large Tables can be partitioned
    - Many partitioning schemes
      - Range partitioning on primary key ~ way of partitioning
- millions of entries
- table split over many machines
- depending on operation, it can be evaluated in parallel



# Select-Project Queries in the Large

there's  
a part  
we  
know  
we  
don't

we only have partial information  
about the string we want to match

① **SELECT \*** all attributes  
**FROM Beers** from table  
**WHERE name like 'Am%'**  
partial match query

② **SELECT name**  
**FROM Beers**  
**WHERE manf = 'Heineken'**

name	manf
A...	Gambrinus
A...	Heineken
...	
B...	Anheuser-Busch

Machine 1

name	manf
C...	MillerCoors
C...	MillerCoors
...	
D...	Duvel Moortgat

Machine 2

name	manf
H...	Heineken
H...	Pabst
...	
H...	Anheuser-Busch

Machine 5

## • Two queries

- ① • Find records for beers whose name starts with 'Am'
- ② • Which beers are made by Heineken?

# Evaluating SP Queries for Large Data

\* do we need to touch all partitions to answer the query? NO!

name is arranged and partitioned based on it

name	manf
A...	Gambrinus
A...	Heineken
...	
B...	Anheuser-Busch

Machine 1

name	manf
C...	MillerCoors
C...	MillerCoors
...	
D...	Duvel Moortgat

Machine 2

name	manf
H...	Heineken
H...	Pabst
...	
H...	Anheuser-Busch

Machine 5

```
SELECT *  
FROM Beers  
WHERE name like 'Am%'
```

evaluation process should only access machine 1 because no other machine will have records for names starting with A

\* so long as system knows partitioning strategy, it can make its job more

efficient

if matters

## • A query processing trick

- Use the partitioning information
  - Just use partition 1!!

# Evaluating SP Queries for Large Data

name	manf
A...	Gambrinus
A...	Heineken
...	
B...	Anheuser-Busch

Machine 1

name	manf
C...	MillerCoors
C...	MillerCoors
...	
D...	Duvel Moortgat

Machine 2

name	manf
H...	Heineken
H...	Pabst
...	
H...	Anheuser-Busch

Machine 5

```
SELECT name  
FROM Beers  
WHERE manf = 'Heineken'
```

can't get away by using partitioning information because partitioning activity is different from attribute on which the query is applied

1. query broadcasted from primary machine to all machines
2. run in local machine
3. results back to primary
4. unioned together
5. return results to client

this will need to go to all partitions

executed in parallel

**Broadcast query**  
In each machine in parallel:  
Select  $manf='Heineken'$  (Beers)  
Project(name)

Gather Partial Results  
Union  
Return  
*(less chance of dealing with large data)*

# Local and Global Indexing

- What if a machine does not have any data for the query attributes?

## Index structures

index — reverse table

- Given value, return records
- Several solutions

- Use local index on each machine
    - main query will go to all machines but lookup will be instant
  - Use a machine index for each value
    - Keeps an account of the machine that contains the record with that value
  - Use a combined index in a global index server
    - ↳ USE BOTH SCHEMES
      - ↳ USES MORE SPACE
      - ↳ QUERIES ARE FASTER
- \* Using an index speeds up query processing

manf	RecordIDs
...	...
MillerCoors	34, 35, 87, 129, ...
Duvel Moortgat	5, 298, 943, 994, ...
Heineken	631, 683, 882, ...
...	...

manf	machinIDs
...	...
MillerCoors	10
Duvel Moortgat	3, 4
Heineken	1, 3, 5

# Pause

QUERIES FOR TWO TABLES

# Querying Two Relations

- Often we need to combine two relations for queries

- Find the beers liked by drinkers who frequent The Great American Bar

- In SQL

```
SELECT DISTINCT beer
  FROM Likes L, Frequent F
 WHERE bar = 'The Great American Bar' AND
       F.drinker = L.drinker
```

conditions

avoids duplicates

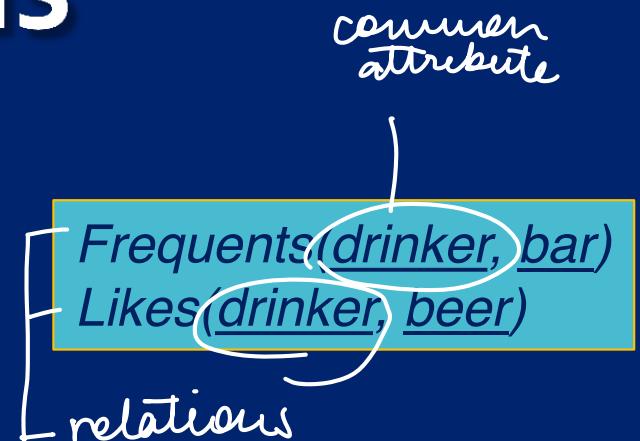
alias L (attrs are unique)

alias F

short cut from L

single table condition

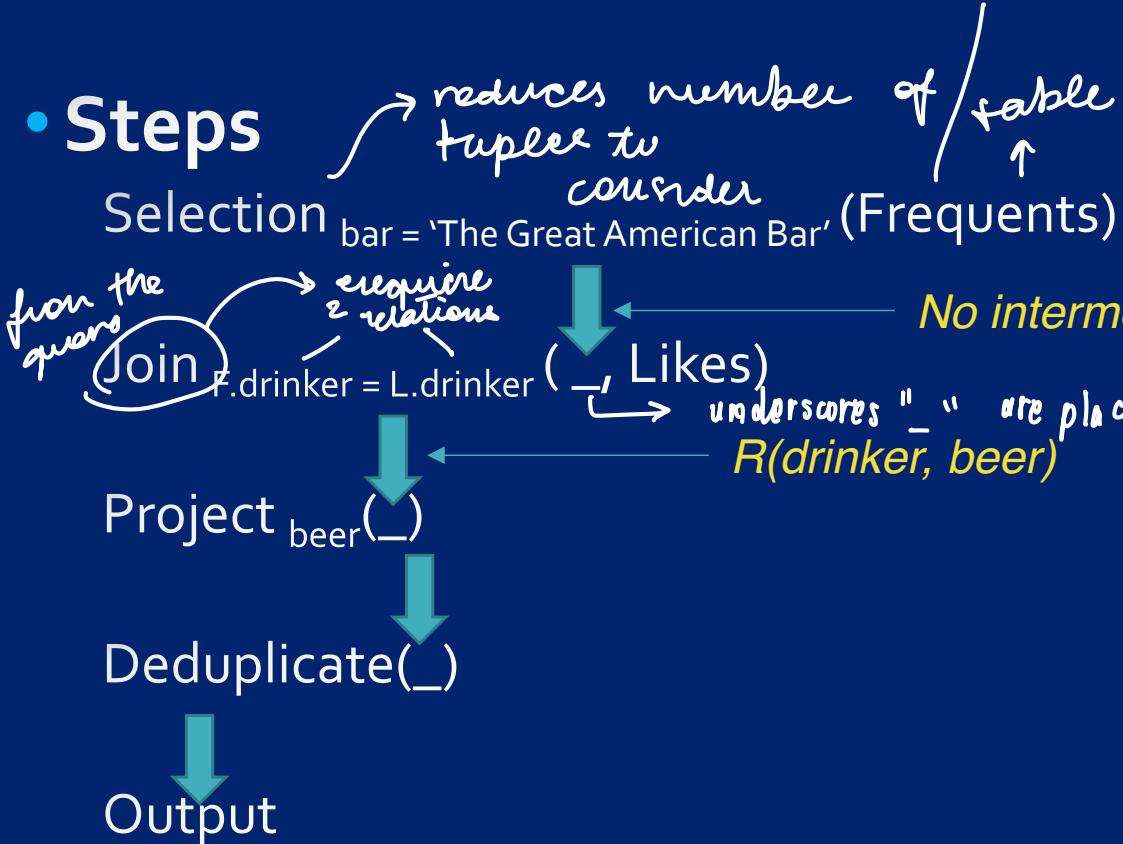
joined condition



# SPJ Queries

- Steps

Selection



*Frequents(drinker, bar)  
Likes(drinker, beer)*

```
SELECT DISTINCT beer  
FROM Likes L, Frequents F  
WHERE bar = 'The Great American Bar'  
AND F.drinker = L.drinker
```

↓  
the selection is piped  
into the join  
operation

# Join in a Distributed Setting

*Frequents(drinker, bar)* → two different  
*Likes(drinker, beer)* → machines

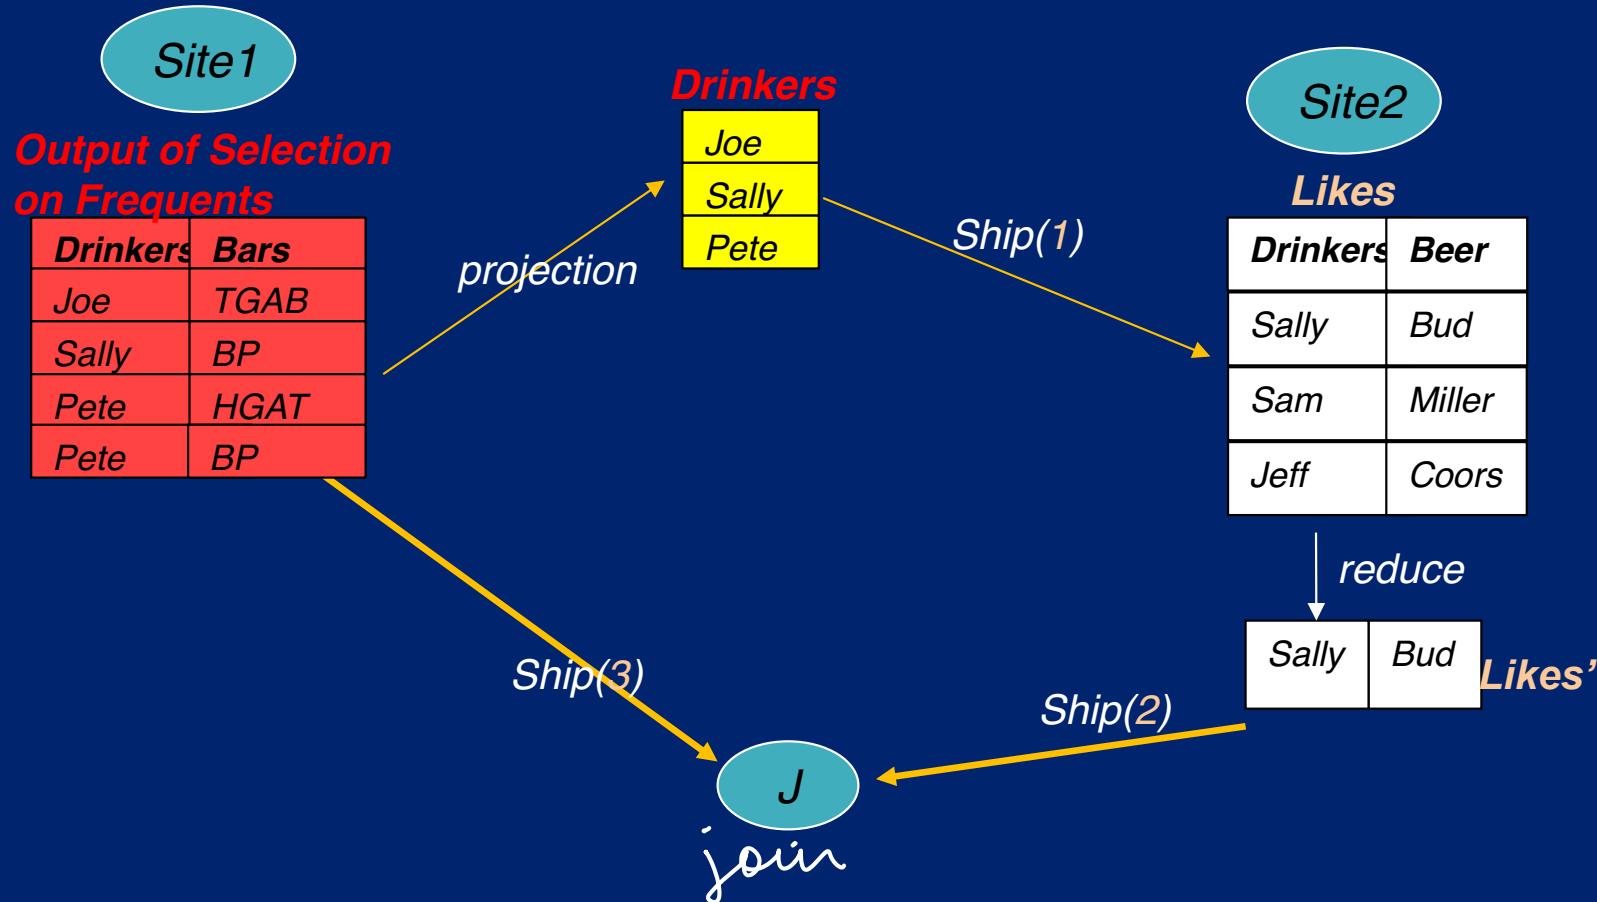
- **Semijoin** operation — move data from one machine to another

- A semijoin from R to S on attribute is used to reduce the data transmission cost
- Computing steps:

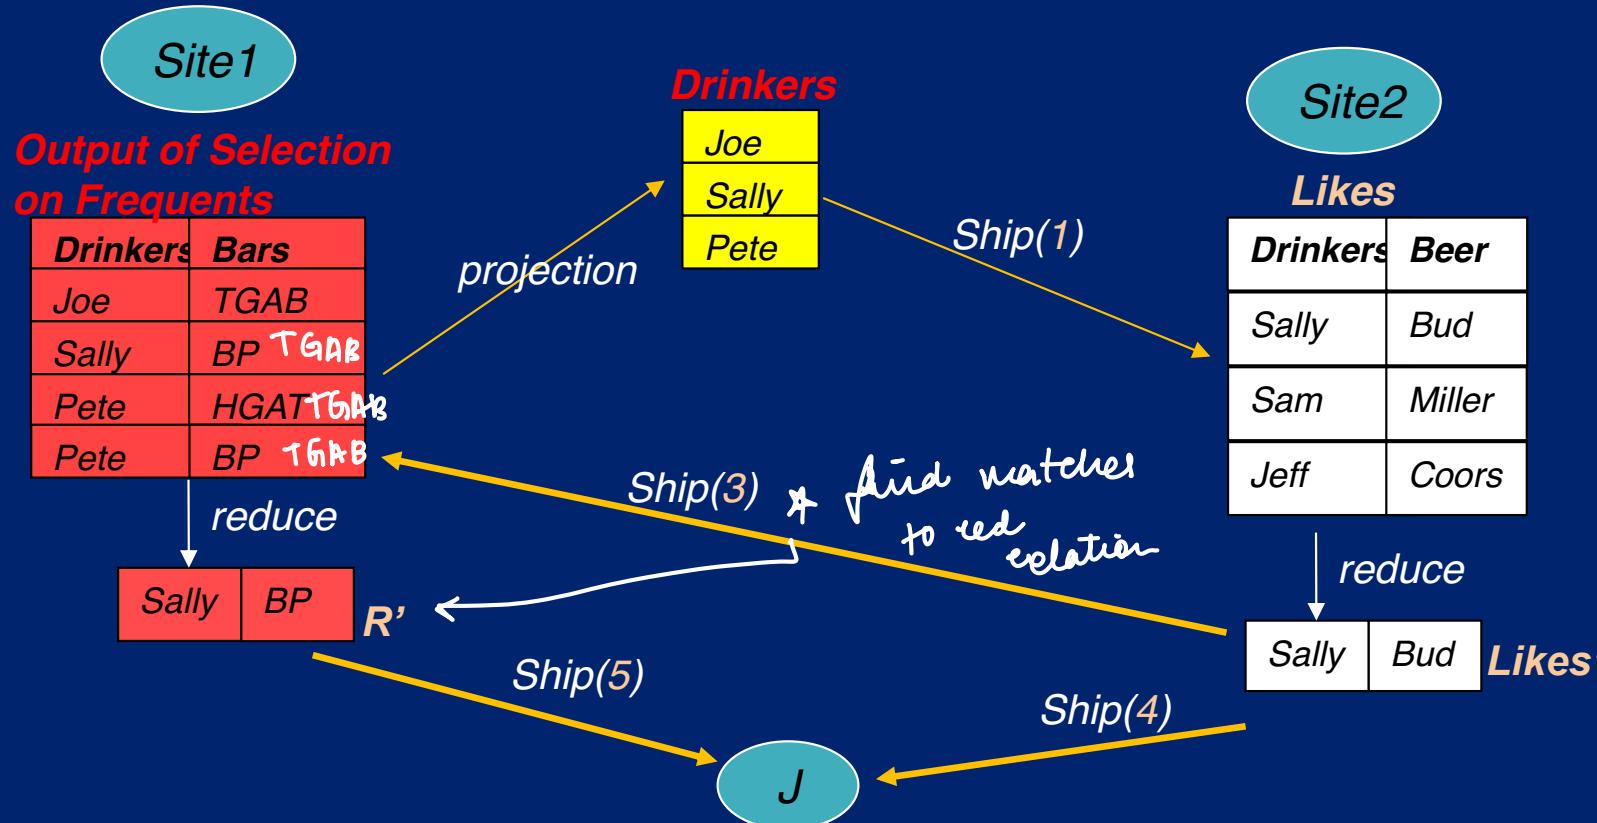
- ① • **Project** R on attribute A and call it (R[A]) – the Drinkers column
- ② • **Ship** this projection ( a semijoin projection) from the site of R to the site of S
- ③ • **Reduce** S to S' by eliminating tuples where attribute A are not matching any value in R[A]

frequents (machine)  
↓  
likes (machine)  
\$ (cost is reduced  
if ship less data)

# Semijoin s: Frequentes—Drinkers → Likes



# Semijoins: Frequentes—Drinkers → Likes



# Pause

# Subqueries

- A slightly complex query
- Find the bars that serve Miller for the same or less price than what TGAB charges for Bud
- We the great american bar may break it into two queries:
  - ① Find the price TGAB charges for Bud
  - ② Find the bars that serve Miller at that price

result is fed as parameter to second query

# Subqueries in SQL

```
SELECT bar  
FROM Sells  
WHERE beer = 'Miller' AND
```

```
    price <= (SELECT price  
              FROM Sells  
             WHERE bar = 'TGAB'  
               AND beer = 'Bud');
```

*The price at  
which TGAB  
sells Bud*

subquery  
(inner  
query)  
• independent  
• uncorrelated

# Subqueries with IN

- Find the name and manufacturer of each beer that Fred does not like

- Query

```
SELECT *  
FROM Beers  
WHERE name
```

NOT IN

```
( SELECT beer  
    FROM Likes  
   WHERE drinker = 'Fred');
```

*Beers(name, manf)*  
*Likes(drinker, beer)*

every name  
does not  
appear here

uncorrelated

# Correlated Subqueries

- Find the name and price of each beer that is more expensive than the average price of beers sold in the bar

```
SELECT beer, price  
FROM Sells s1  
WHERE price >  
(SELECT AVG(price)  
FROM Sells s2  
WHERE s1.bar = s2.bar)
```

① average of every bar  
② compare price of each beer to average

stores average once computed and reuses it

Bar	Beer	Price
HGAT	Bud	5
BP	Michelob	4
TGAB	Heineken	6
HGAT	Guinness	10

for every table processed by the outer query,  
one needs to compute the inner query for that bar  
↓  
correlated

# Aggregate Queries

5  
3  
4  
— not repeated

- **Example**

- Find the average price of Bud:
- SELECT AVG(price)
- FROM Sells
- WHERE beer = 'Bud';

5  
3  
4  
4  
5

→ 4. 4.2a

SELECT AVG (DISTINCT price)  
FROM Sells  
WHERE beer = 'Bud'

4

- **Other aggregate functions**

- SUM, MIN, MAX, COUNT, ...

# GROUP BY Queries

- Find for each drinker the average price of Bud at the bars they frequent

SELECT drinker, AVG(price)

FROM Frequent, Sells  
WHERE beer = 'Bud' AND  
Frequent.bar = Sells.bar  
GROUP BY drinker

Drinker	Bar	Price
Pete	HGAT	5
Pete	BP	4
Joe	TGAB	6
Joe	HGAT	5

↓

Drinker	Price
Pete	4.5
Joe	5.5

join attribute : statistical aggregate by groups  
grouping variable : creates one result row for each drinker and place average price

# Grouping Aggregates over Partitioned Data

Drinker	Bar	Price
Pete	HGAT	5
Pete	BP	4
Joe	TGAB	6
John	HGAT	5

① data gets repartitioned by grouping attribute → (drinker)

Drinker	Bar	Price
Pete	HGAT	5
Pete	BP	4
Pete	BO	6
Joe	TGAB	6

result of selection in two different machines

Drinker	Bar	Price
Pete	BO	6
John	BP	4
Sally	TGAB	6
Sally	HGAT	5

each machine groups its own data locally

② aggregate function computed locally

Drinker	Price
Pete	5
Joe	6

Drinker	Bar	Price
John	HGAT	5
John	BP	4
Sally	TGAB	6
Sally	HGAT	5

Drinker	Price
John	4.5
Sally	5.5

\* there are variants