


Querying Relational Data with Postgres

By the end of this activity, you will be able to:

1. View table and column definitions, and perform SQL queries in the Postgres shell
2. Query the contents of SQL tables
3. Filter table rows and columns
4. Combine two tables by joining on a column

Step 1. **Open a terminal window and start Postgres shell.** Open a terminal window by clicking on the square black box on the top left of the screen.



Next, start the Postgres shell by running *psql*:

```
[cloudera@quickstart big-data-3]$ psql
psql (8.4.20)
Type "help" for help.

cloudera=#
```

Step 2. **View table and column definitions.** We can list the tables in the database with the *\d* command:

```
cloudera=# \d
List of relations
Schema | Name          | Type  | Owner
-----+-----+-----+-----
public | adclicks      | table | cloudera
public | buyclicks     | table | cloudera
public | gameclicks    | table | cloudera
(3 rows)
```

The database contains three tables: *adclicks*, *buyclicks*, and *gameclicks*. We can see the column definitions of the *buyclicks* table by running `\d buyclicks`:

```
cloudera=# \d buyclicks
Table "public.buyclicks"
  Column          |          Type          | Modifiers
-----+-----+-----
 timestamp       | timestamp without time zone | not null
 txid            | integer                | not null
 usersessionid   | integer                | not null
 team            | integer                | not null
 userid          | integer                | not null
 buyid           | integer                | not null
 price           | double precision       | not null
```

This shows that the *buyclicks* table has seven columns, and what each column name and data type is.

Step 3. Query table. We can run the following command to view the contents of the *buyclicks* table:

```
select * from buyclicks;
```

The *select ** means we want to query all the columns, and *from buyclicks* denotes which table to query. Note that all query commands in the Postgres shell must end with a semi-colon.

The result of the query is:

timestamp	txid	usersessionid	team	userid	buyid	price
2016-05-26 15:36:54	6004	5820	9	1300	2	3
2016-05-26 15:36:54	6005	5775	35	868	4	10
2016-05-26 15:36:54	6006	5679	97	819	5	20
2016-05-26 16:36:54	6067	5665	18	121	2	3
2016-05-26 17:06:54	6093	5709	11	2222	5	20
2016-05-26 17:06:54	6094	5798	77	1304	5	20
2016-05-26 18:06:54	6155	5920	9	1027	5	20
2016-05-26 18:06:54	6156	5697	35	2199	2	3
2016-05-26 18:36:54	6183	5893	64	1544	5	20
2016-05-26 18:36:54	6184	5697	35	2199	1	2
2016-05-26 19:36:54	6243	5659	13	1623	4	10

You can hit <space> to scroll down, and *q* to quit.

Step 4. **Filter rows and columns.** We can query only the *price* and *userid* columns with the following command:

```
select price, userid from buyclicks;
```

The result of this query is:

price	userid
3	1300
10	868
20	819
3	121
20	2222
20	1304
20	1027
3	2199
20	1544

We can also query rows that match a specific criteria. For example, the following command queries only rows with a price greater than 10:

```
select price, userid from buyclicks where price > 10;
```

The result is:

price	userid
20	819
20	2222
20	1304
20	1027
20	1544
20	1065
20	2221

Step 5. **Perform aggregate operations.** The SQL language provides many aggregate operations. We can calculate the average price:

```
cloudera=# select avg(price) from buyclicks;
          avg
-----
 7.26399728537496
(1 row)
```

We can also calculate the total price:

```
cloudera=# select sum(price) from buyclicks;
          sum
-----
      21407
(1 row)
```

The complete list of aggregate functions for Postgres 8.4 (the version installed on the Cloudera VM) can be found here: <https://www.postgresql.org/docs/8.4/static/functions-aggregate.html>

Step 6. **Combine two tables.** We combine the contents of two tables by matching or joining on a single column. If we look at the definition of the *adclicks* table:

```
cloudera=# \d adclicks
```

Column	Type	Modifiers
timestamp	timestamp without time zone	not null
txid	integer	not null
usersessionid	integer	not null
teamid	integer	not null
userid	integer	not null
adid	integer	not null
adcategory	character varying(11)	not null

We see that *adclicks* also has a column named *userid*. The following query combines the *adclicks* and *buyclicks* tables on the *userid* column in both tables:

```
from adclicks join buyclicks on adclicks.userid = buyclicks.userid;

select adid, buyid, adclicks.userid
```

This query shows the columns *adid* and *userid* from the *adclicks* table, and the *buyid* column from the *buyclicks* table. The *from adclicks join buyclicks* denotes that we want to combine these two tables, and *on adclicks.userid = buyclicks.userid* denotes which two columns to use when the tables are combined.

The result of the query is:

adid	buyid	userid
2	5	611
2	4	611
2	4	611
2	5	611
2	4	611
2	1	611
21	1	1874
21	1	1874
21	3	1874
21	1	1874
21	2	1874

Enter `\q` to quit the Postgres shell.