Introduction to Apache Spark

After this video you will be able to..

- List the main motivations for the development of Spark
- Draw the Spark stack as a layer diagram
- Explain the functionality of the components in the Spark stack

Why Spark?

Hadoop MapReduce Shortcomings

restrict pipelines

- Only for Map and Reduce based computations
- Relies on reading data from HDFS
- Native support for Java only
- No interactive shell support
- No support for streaming

Spark was made to overcome this shortcomings and provide an expressive cluster computing environment

- interactive querying
- efficient iterative analytics
- streaming data processing

Basics of Data Analysis with Spark

- Expressive programing model
- In-memory processing
- Support for diverse workloads
- Interactive shell

The Spark Stack

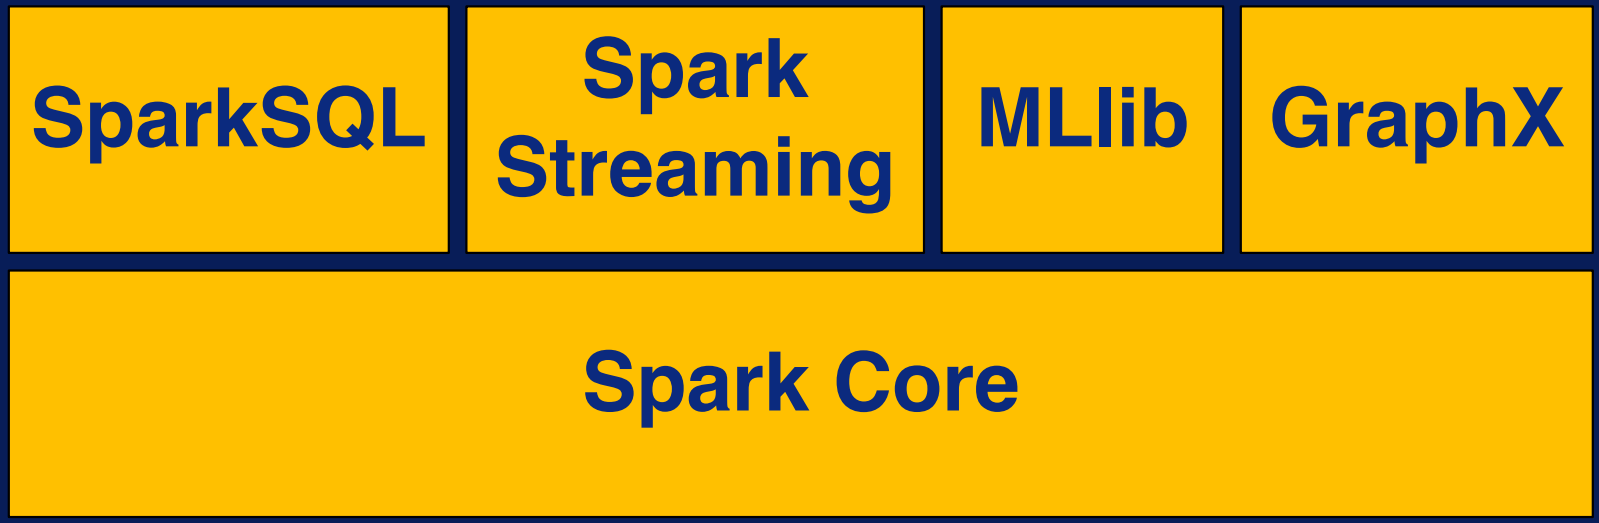

The Spark Stack

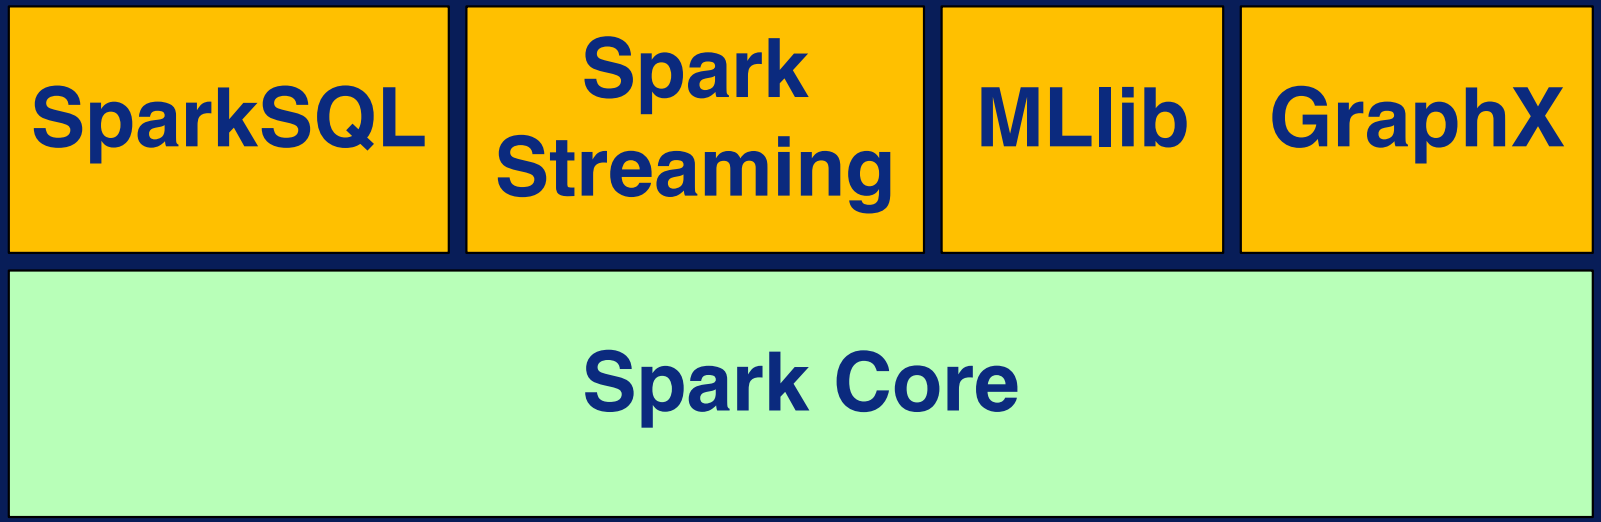

- where the core capabilities of the spark framework are implemented
  - support for distributed scheduling
  - memory management
  - fault tolerance
- interaction with schedulers
- APIs for resilient distributed datasets (RDDs) — main programming abstraction

The Spark Stack

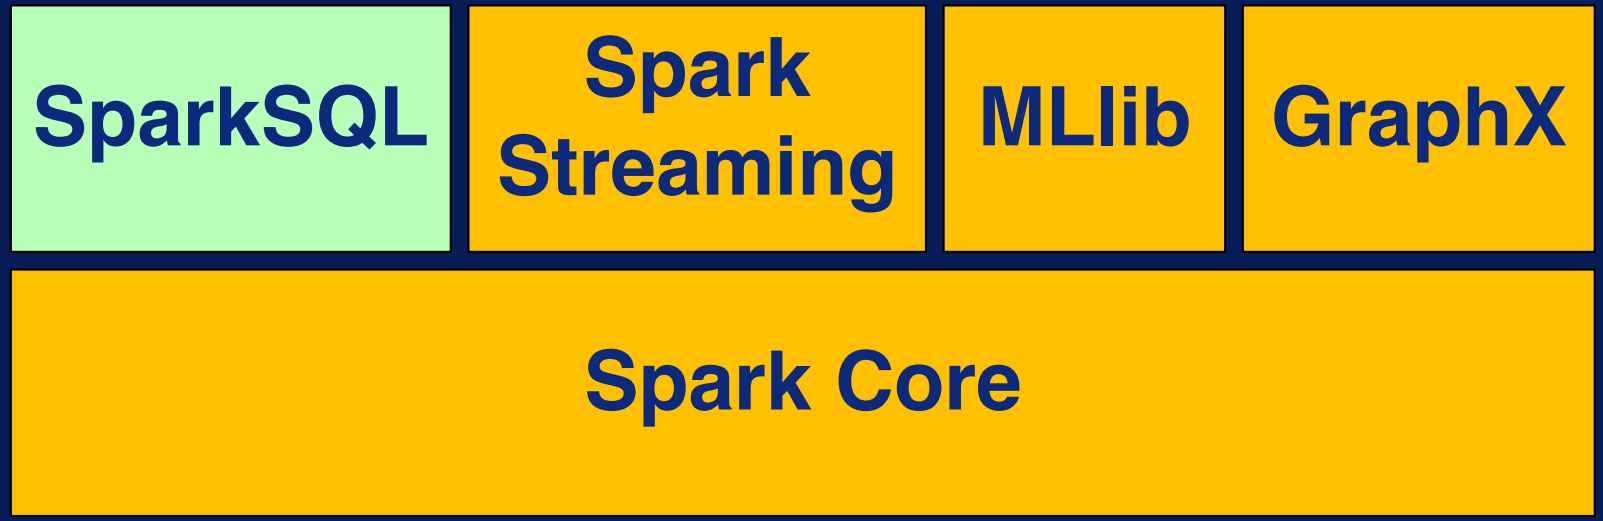

- provides querying structured and unstructured data through a common query languages
  - connects to many data sources
  - APIs to convert query results to RDDs

The Spark Stack

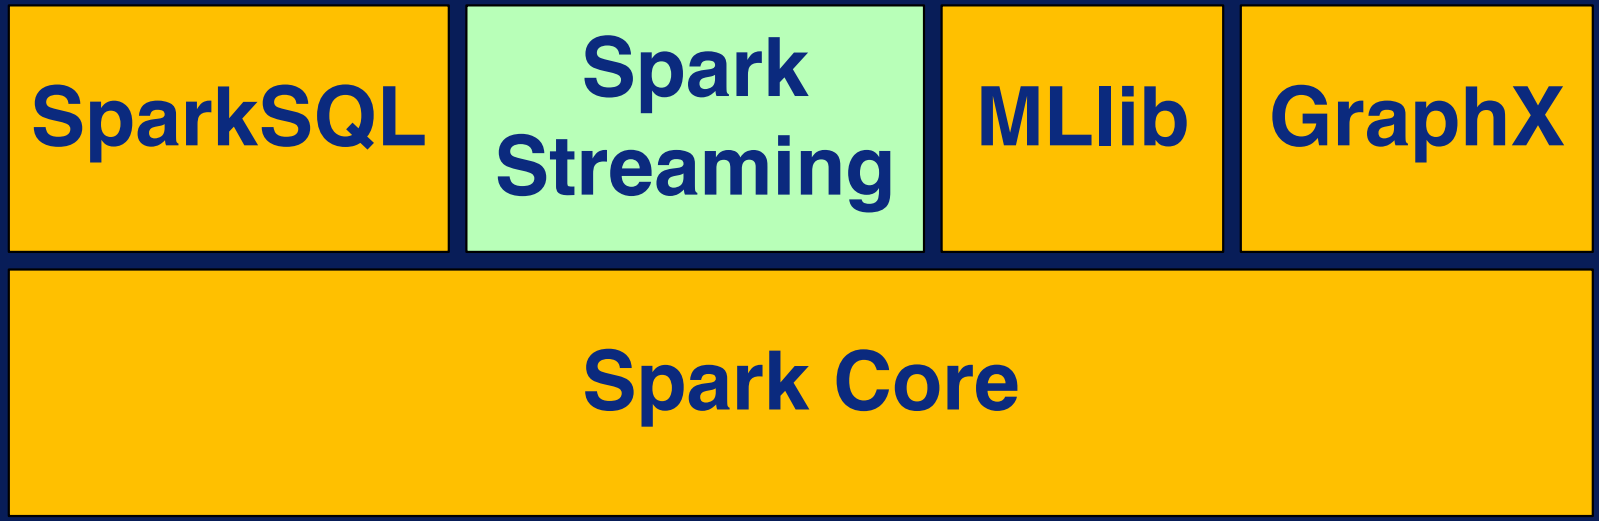

- for streaming data manipulations
- enables creating small aggregates of data incoming from streaming data ingestion systems (micro batches — aggregated datasets) ⇝ can be converted to RDD

The Spark Stack

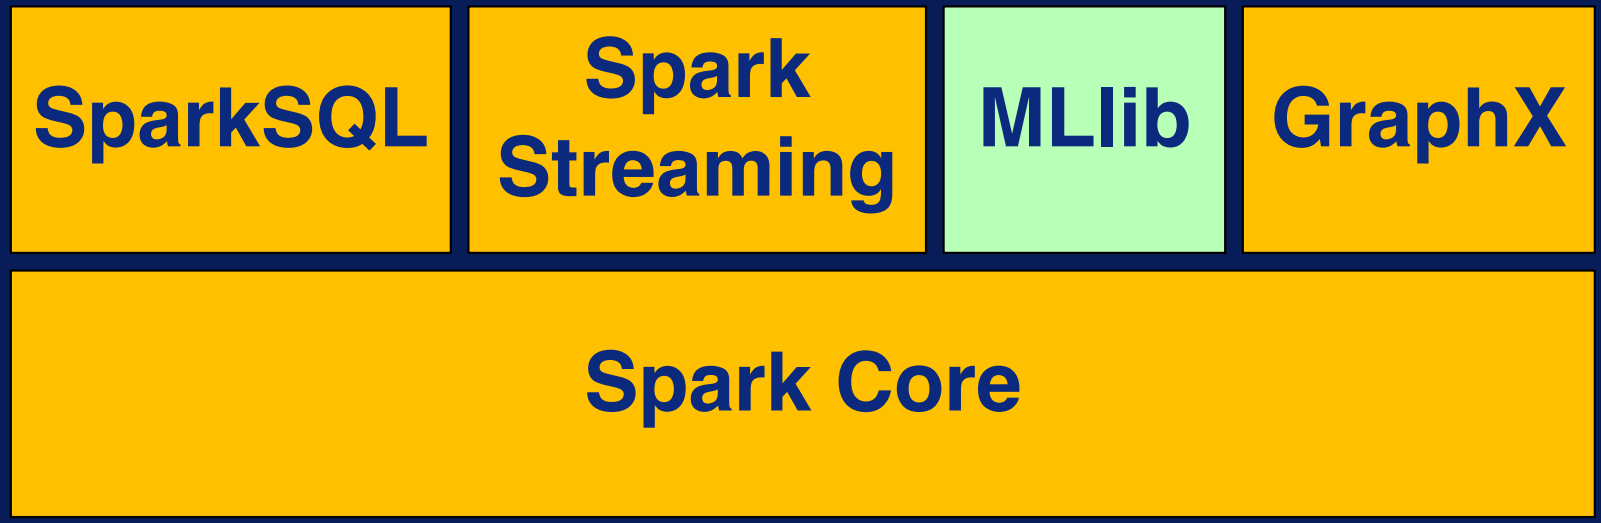

- spark's native library for machine learning algorithms as well as model evaluation
- designed to scale out using spark

The Spark Stack

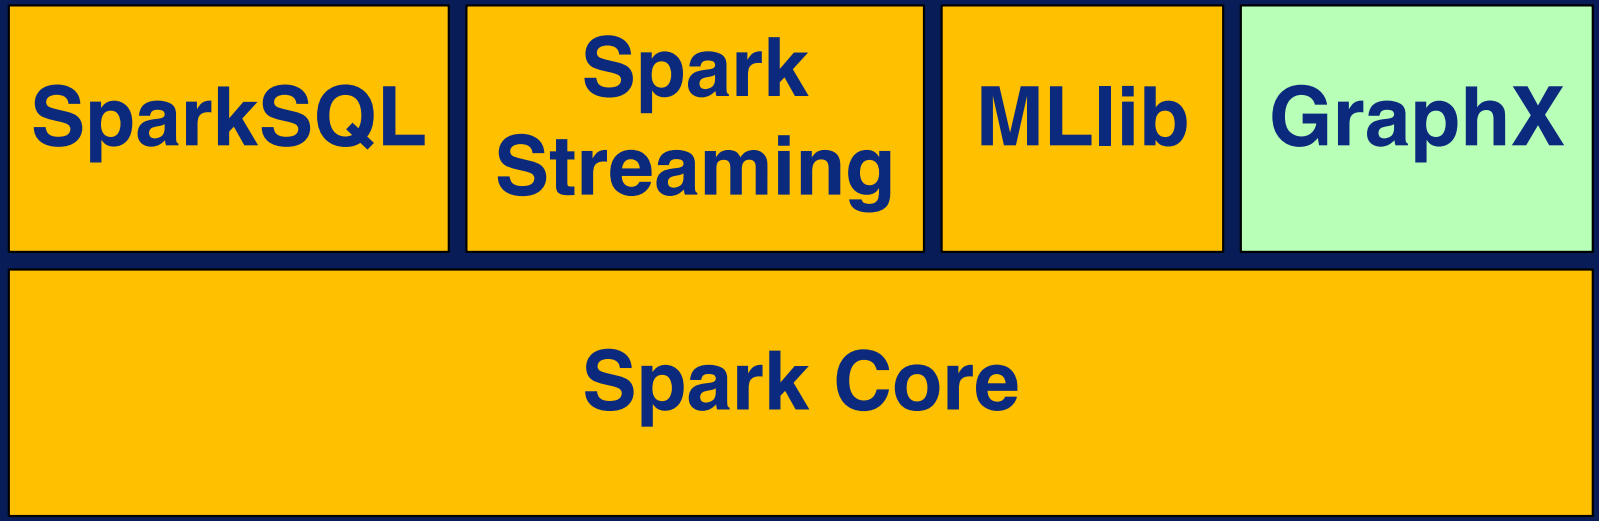

- graph analytics library
- enables the vertex edge data model of graphs to be converted into RDDs
+ scalable implementations of a graph processing algorithms

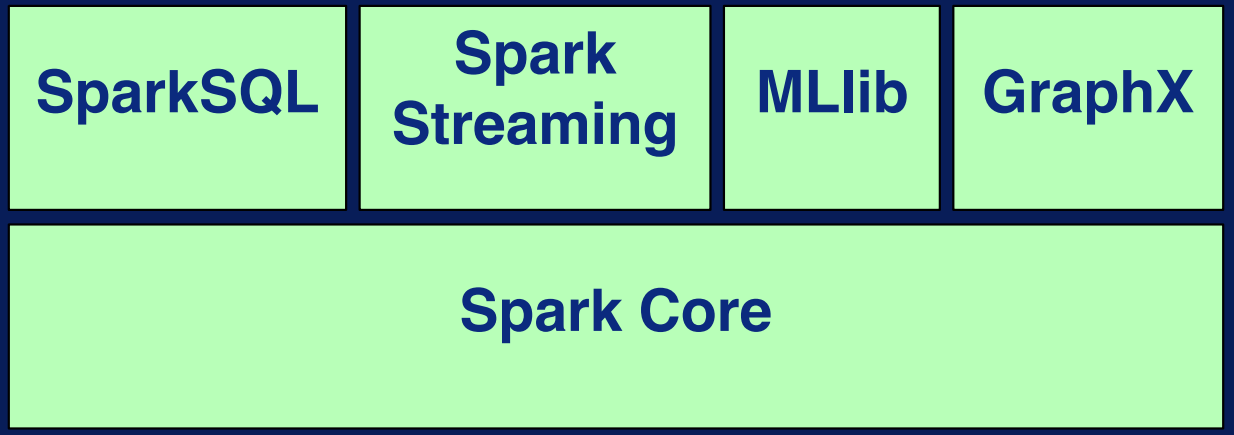

Explore — conduct exploratory analysis

→ * through these layers spark provides ↑ diverse, scalable, interactive — manage, wrangle, and analyze (of big data)

Build — + create data pipelines

Scale — + enabling scaling across commodity computing clusters & cloud environments