Getting Started with Spark: The Architecture and Basic Concepts

After this video you will be able to..

- Describe how Spark does in-memory processing using the RDD abstraction
- Explain the inner workings of the Spark architecture
- Summarize how Spark manages and executes code on Clusters

What does in memory processing mean?

MapReduce

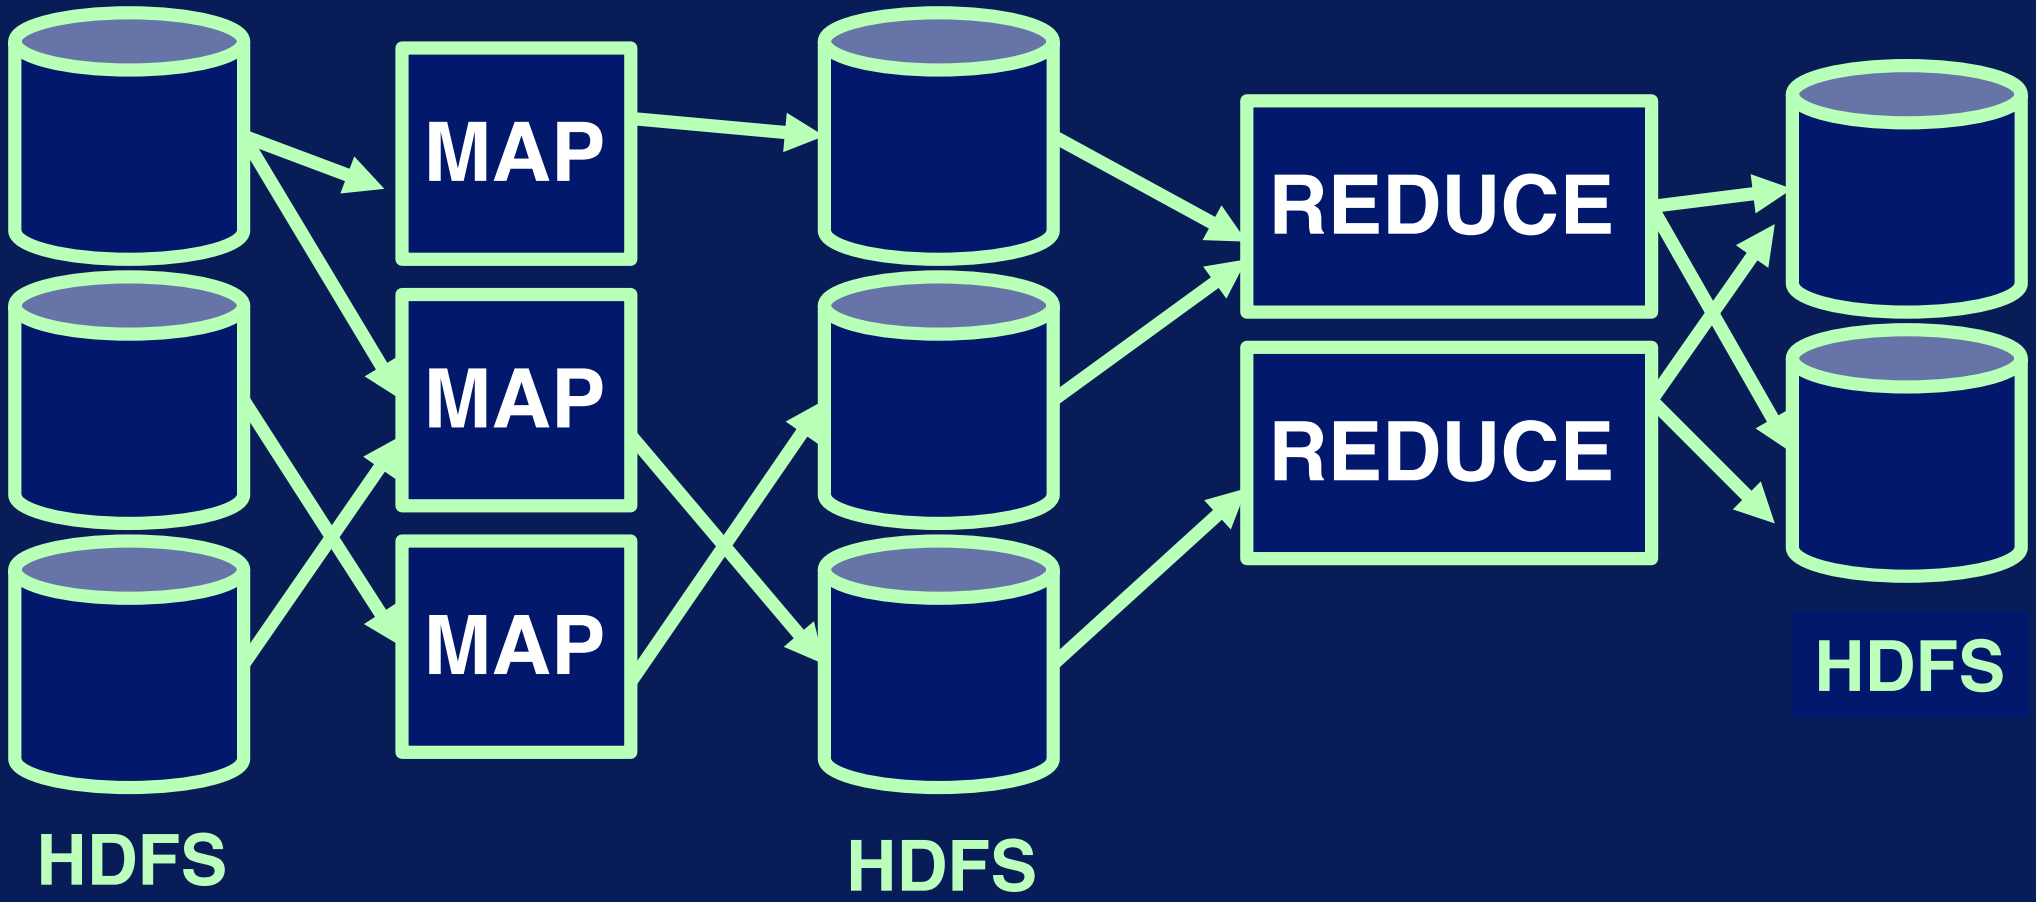

Spark

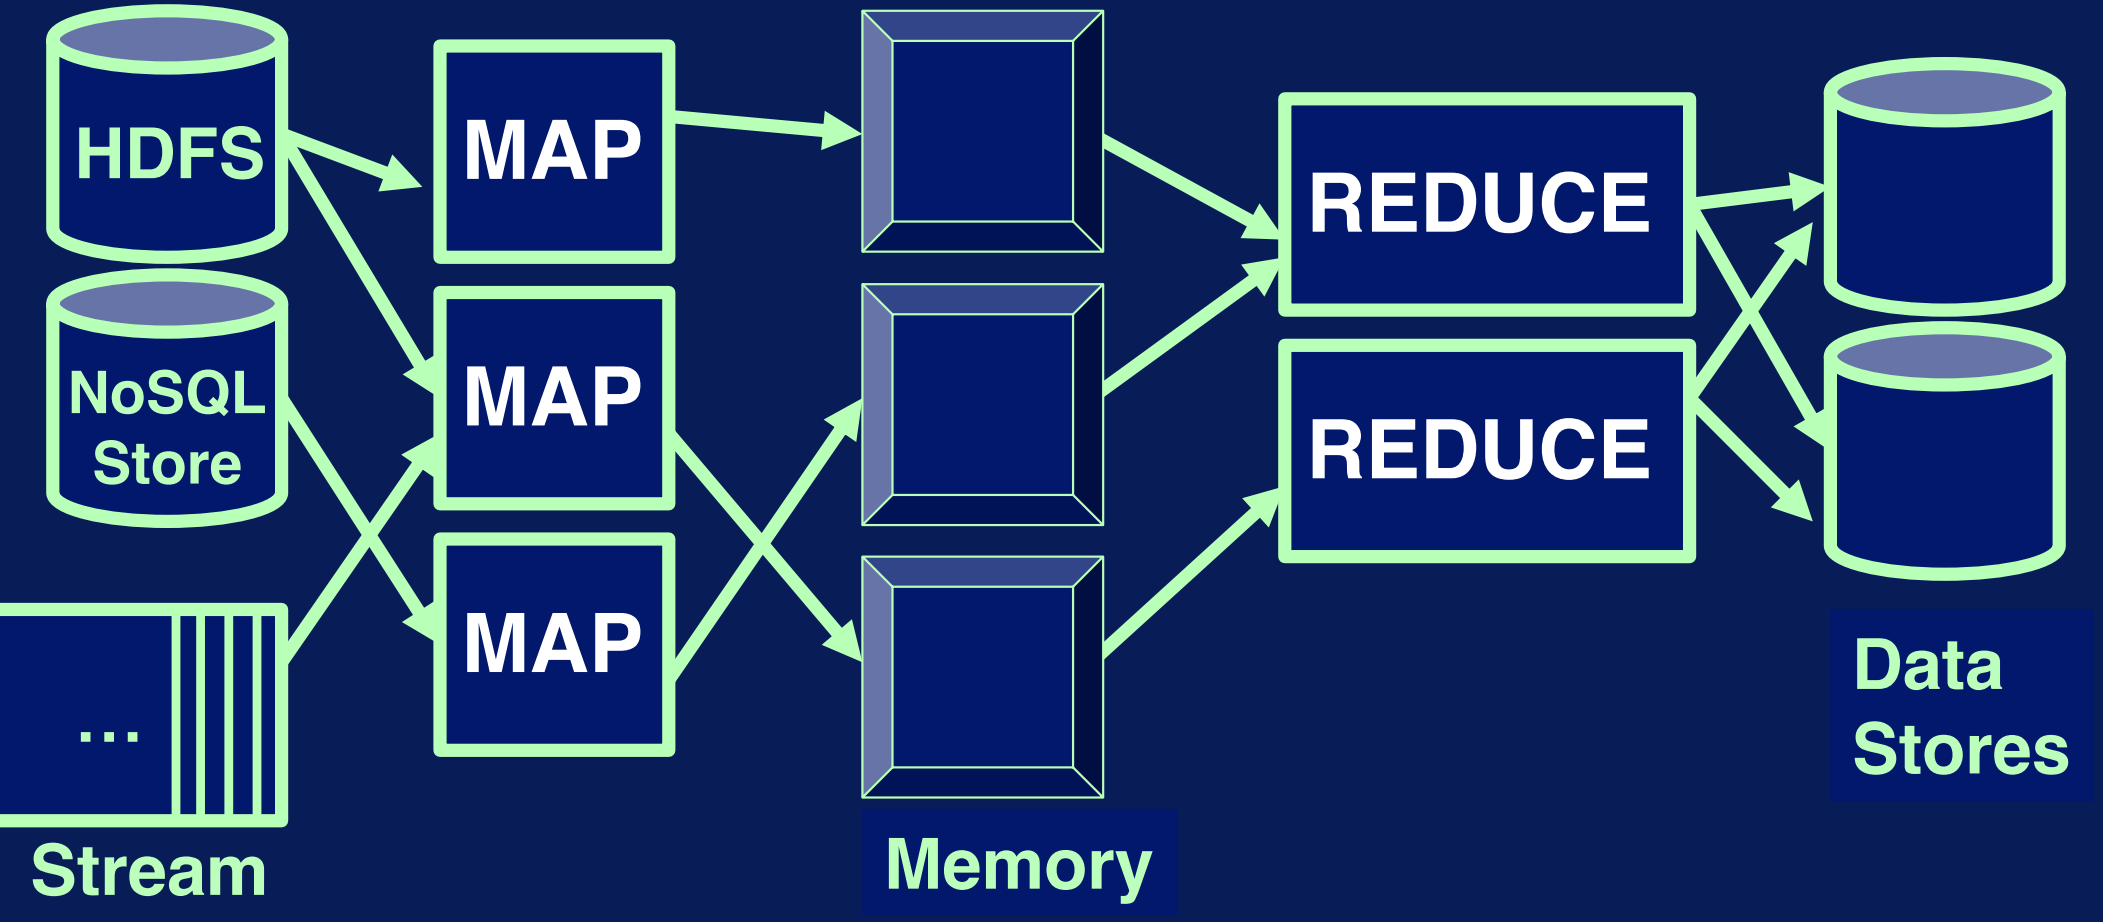
HDFS
NoSQL Store
...
Stream
MAP
MAP
MAP
Memory
REDUCE
REDUCE
Data Stores

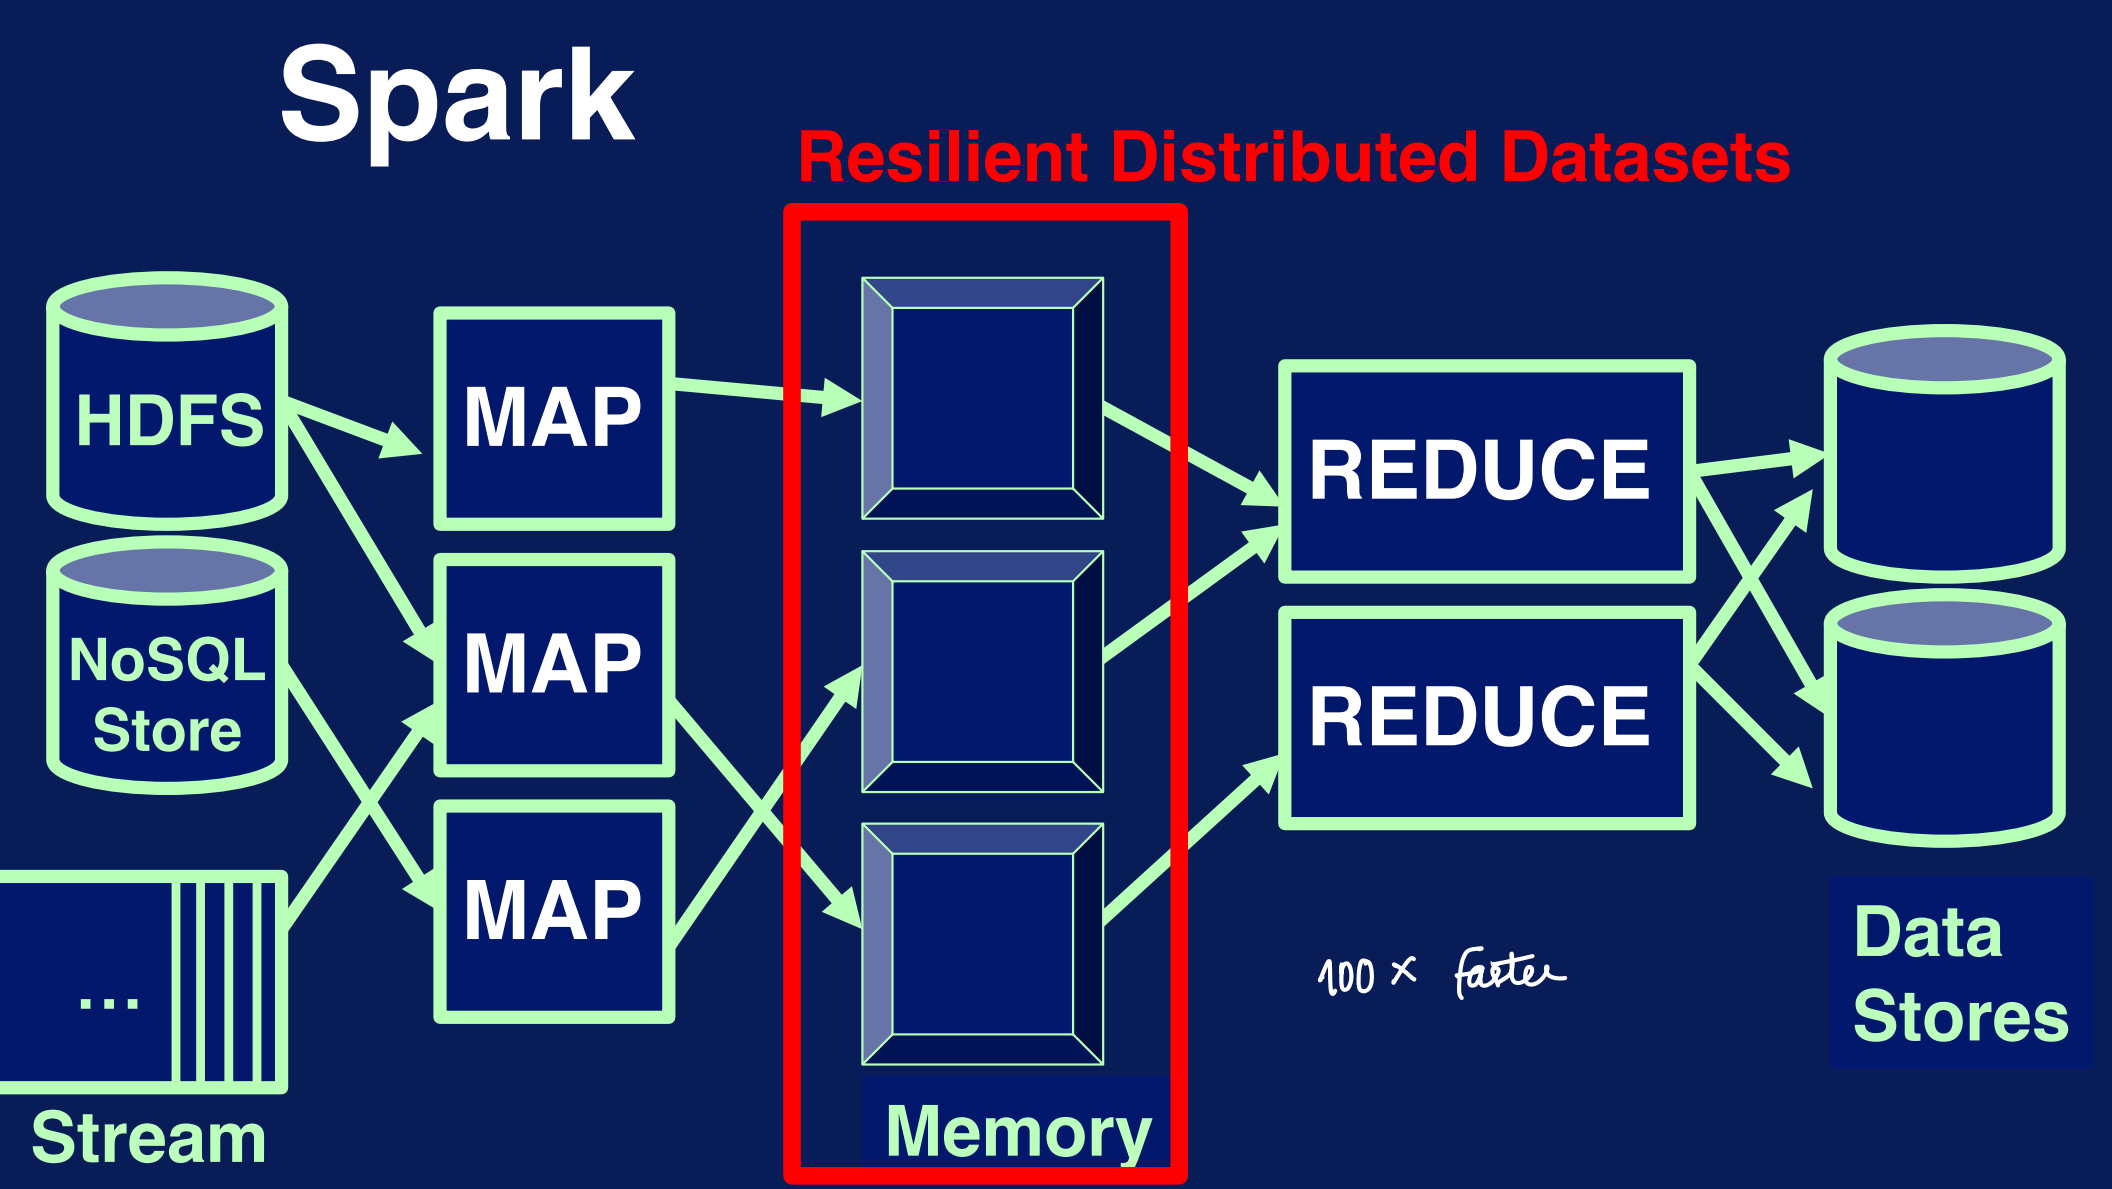
Spark
Resilient Distributed Datasets
HDFS
NoSQL Store
...
Stream
MAP
MAP
MAP
Memory
REDUCE
REDUCE
100 × faster
Data Stores

Resilient Distributed Datasets

Dataset

Data storage created from:
HDFS, S3, HBase, JSON, text,
Local hierarchy of folders

Or created transforming another RDD

Resilient Distributed Datasets

Distributed

Distributed across the cluster of machines

Divided in partitions, atomic chunks of data

Resilient Distributed Datasets

Resilient

Recover from errors, e.g. node failure, slow processes

Track history of each partition, re-run

* it is common to have node failures

↓

its important to recover without losing work already done

spark knows which are the partitions needed to recreate the partition in case it gets lost

Spark Architecture

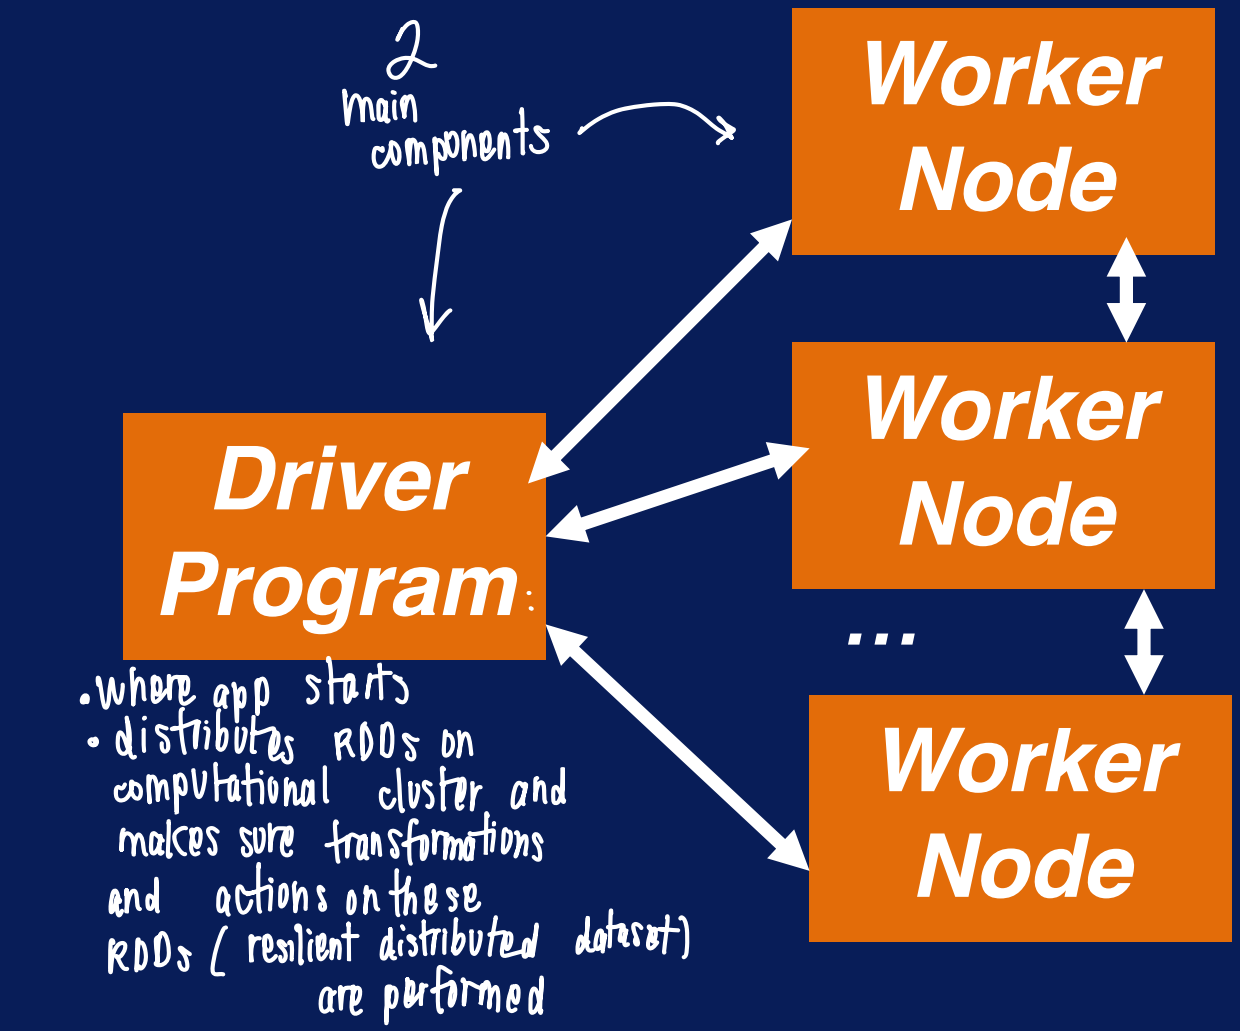
2
main
components
Worker
Node
Worker
Node
...
Worker
Node
Driver
Program:
• where app starts
• distributes RDDs on computational cluster and makes sure transformations and actions on these RDDs (resilient distributed dataset) are performed

Driver Program

```
In [1]: lines = sc.textFile("hdfs:/user/cloudera/words.txt")
```

- creates a connection to a Spark cluster / or local Spark through a Spark context object
  - default: sc (spark context)

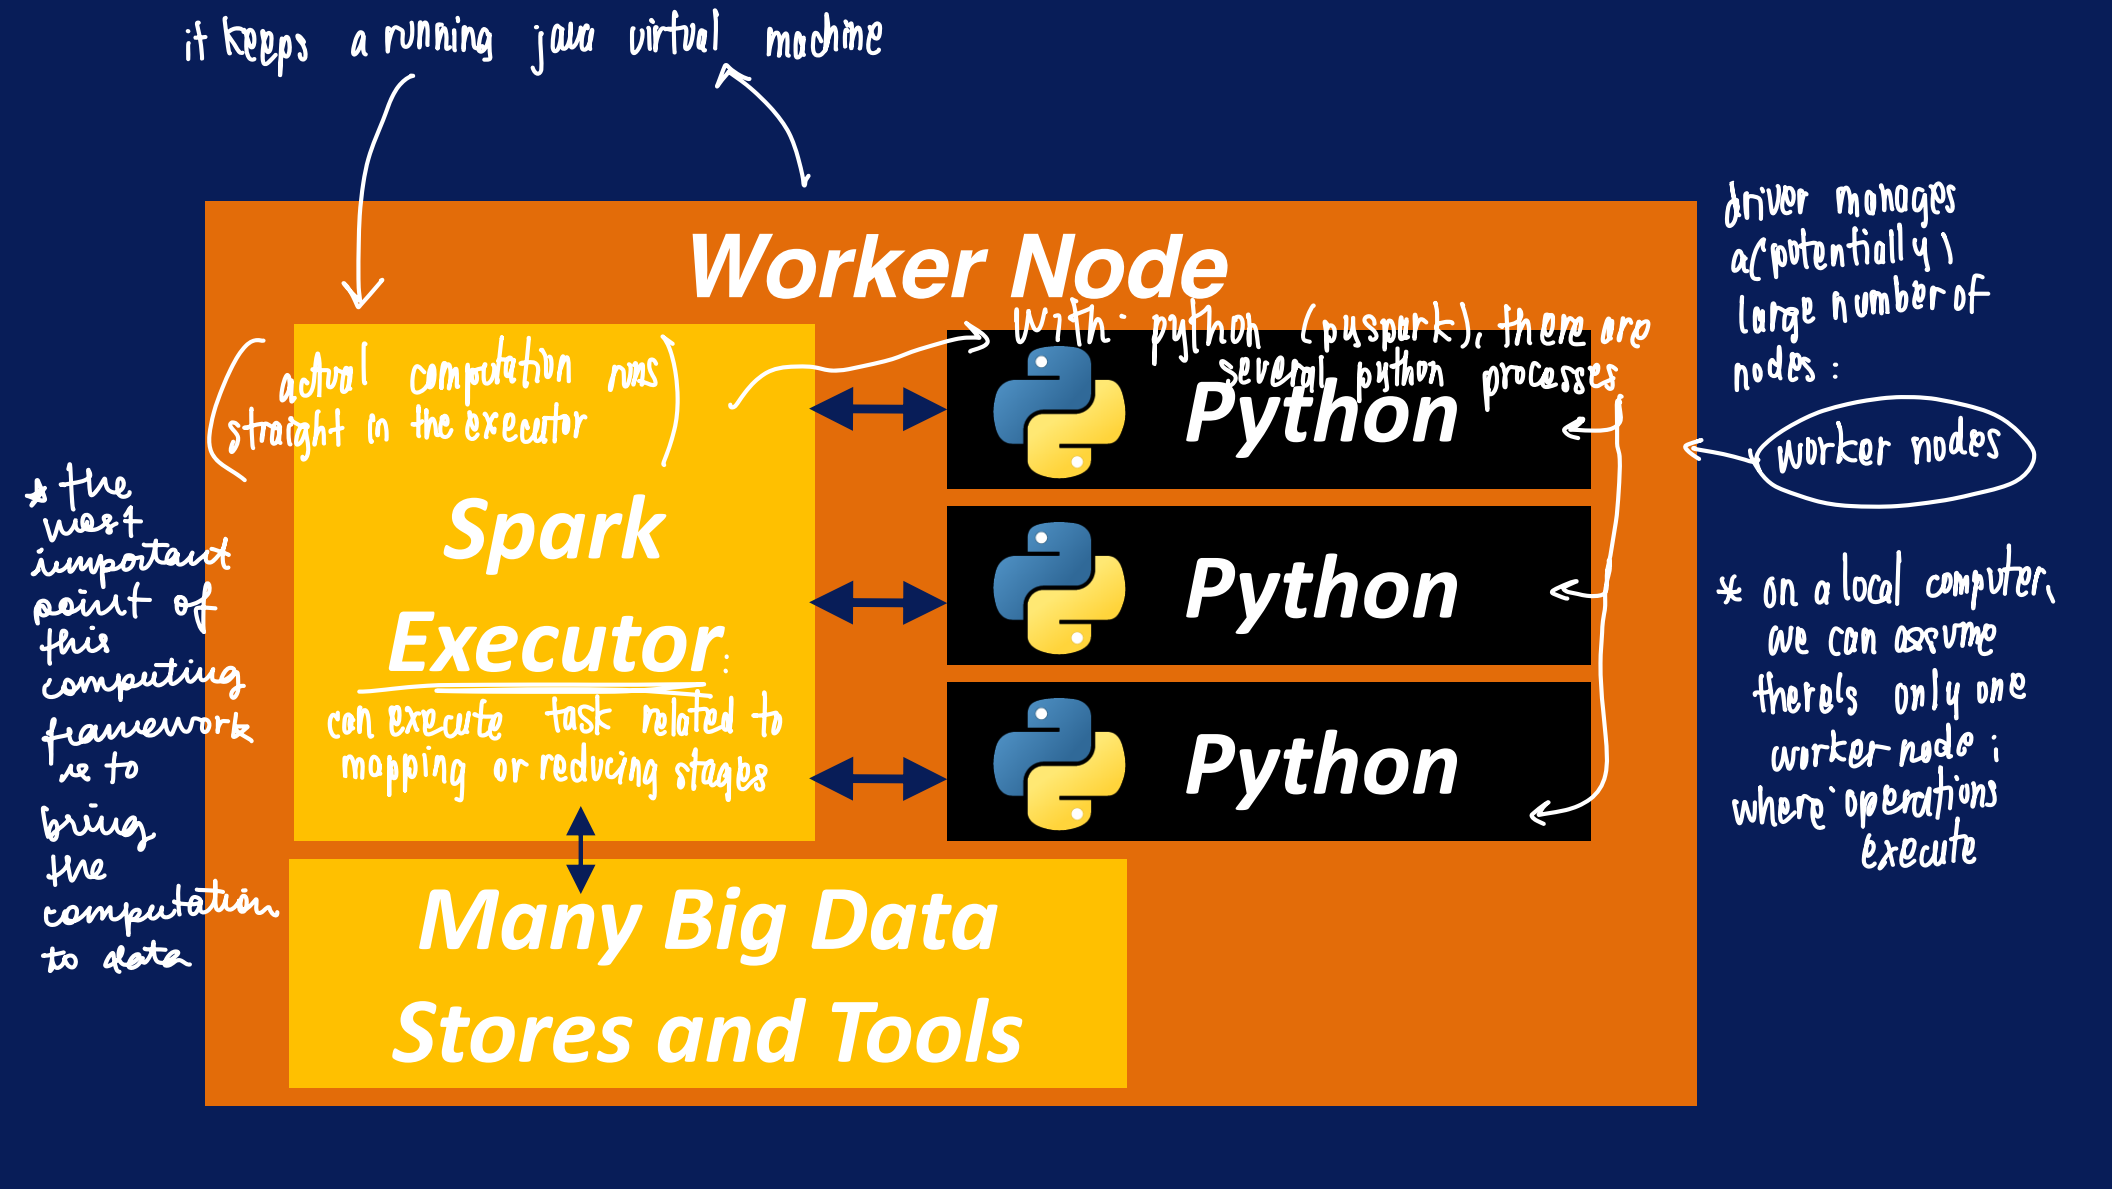
it keeps a running java virtual machine
Worker Node
actual computation runs straight in the executor
Spark Executor:
can execute task related to mapping or reducing stages
With: python (pyspark), there are several python processes
Python
Python
Python
Many Big Data Stores and Tools
driver manages a(potentially) large number of nodes:
worker nodes
* the most important point of this computing framework is to bring the computation to data
* on a local computer, we can assume there is only one worker node: where operations execute

* it is important to have a system that can automatically manage provisioning and restarting of these nodes

many **Worker Nodes**

running tasks internally

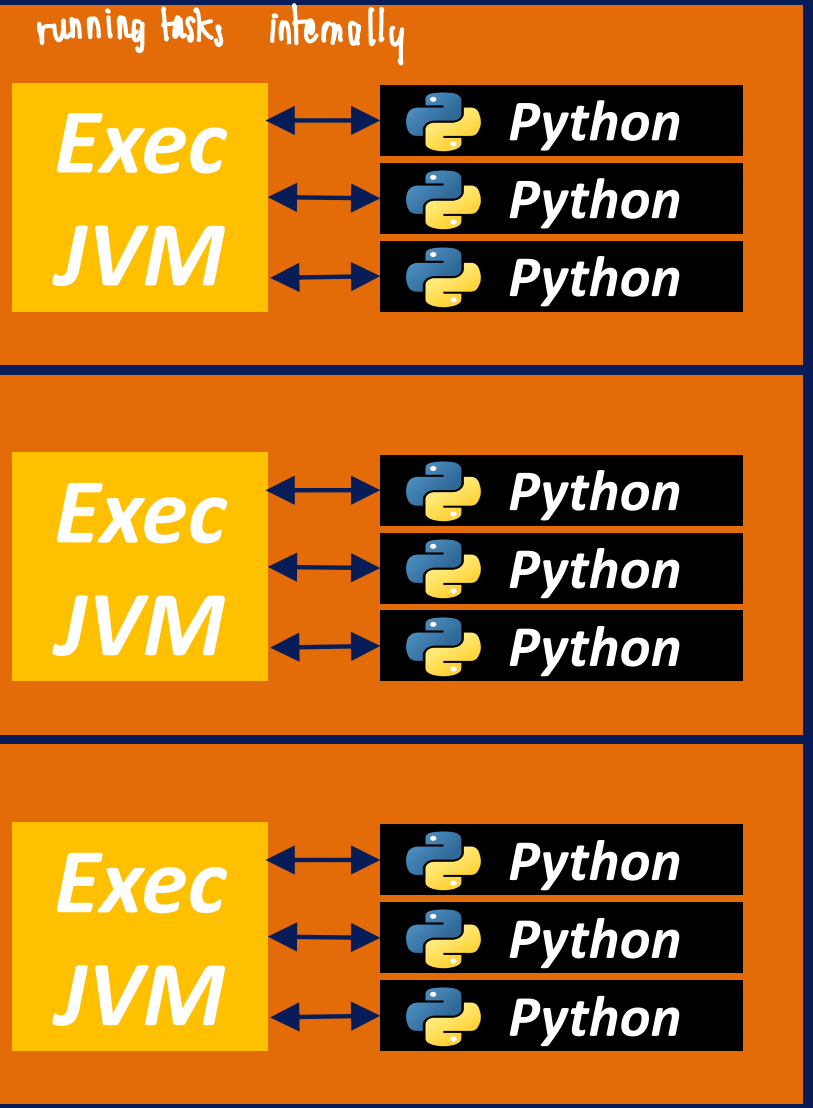

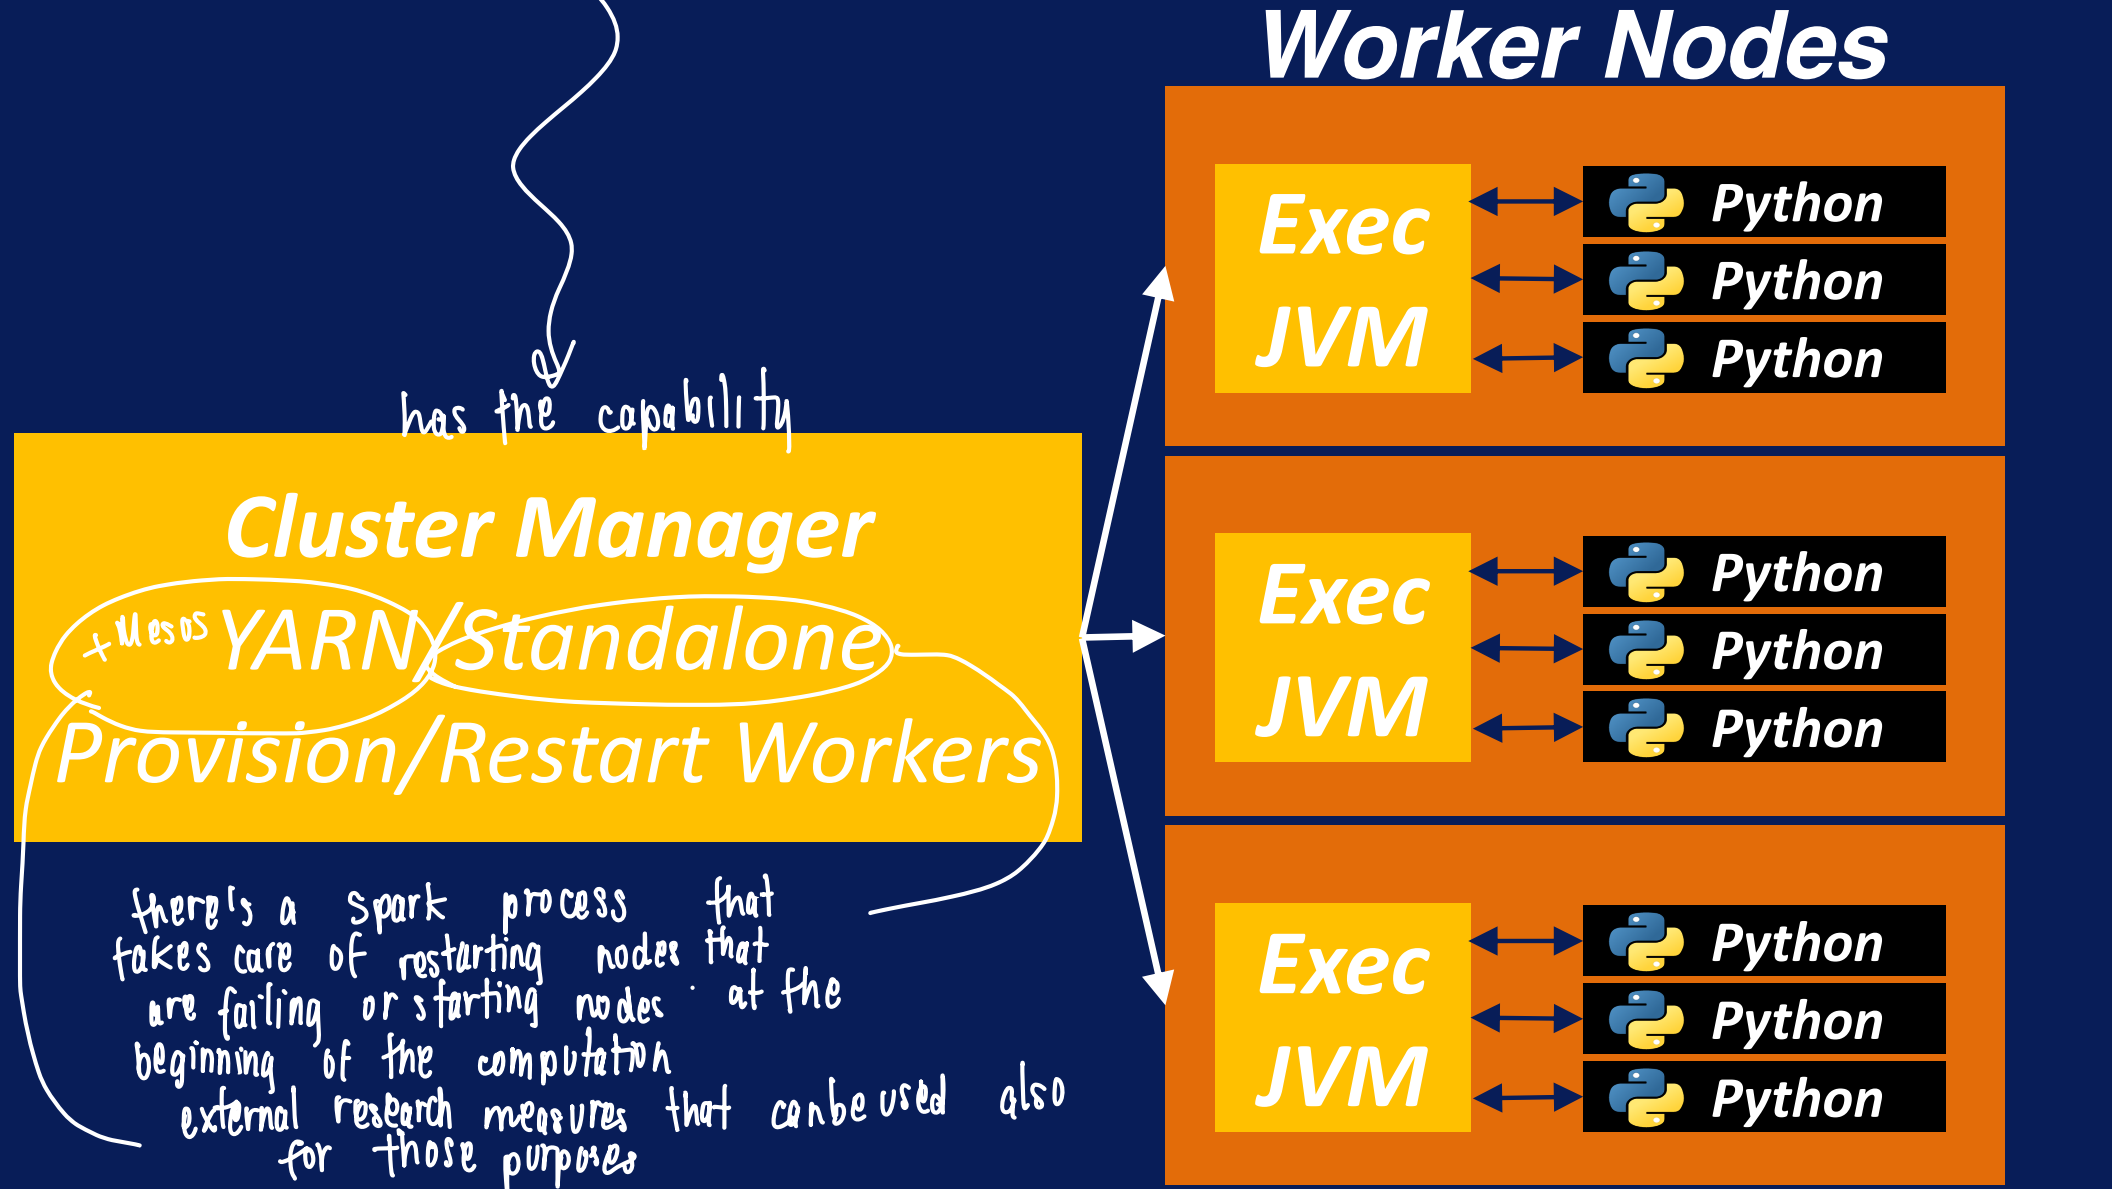
has the capability
Cluster Manager
+Mesos YARN/Standalone
Provision/Restart Workers
there's a spark process that takes care of restarting nodes that are failing or starting nodes at the beginning of the computation
external research measures that can be used also for those purposes
Worker Nodes
Exec JVM
Python
Python
Python
Exec JVM
Python
Python
Python
Exec JVM
Python
Python
Python

Which cluster manager?

* choosing

how to pick the right cluster manager for your organization:

http://www.agildata.com/apache-spark-cluster-managers-yarn-mesos-or-standalone/

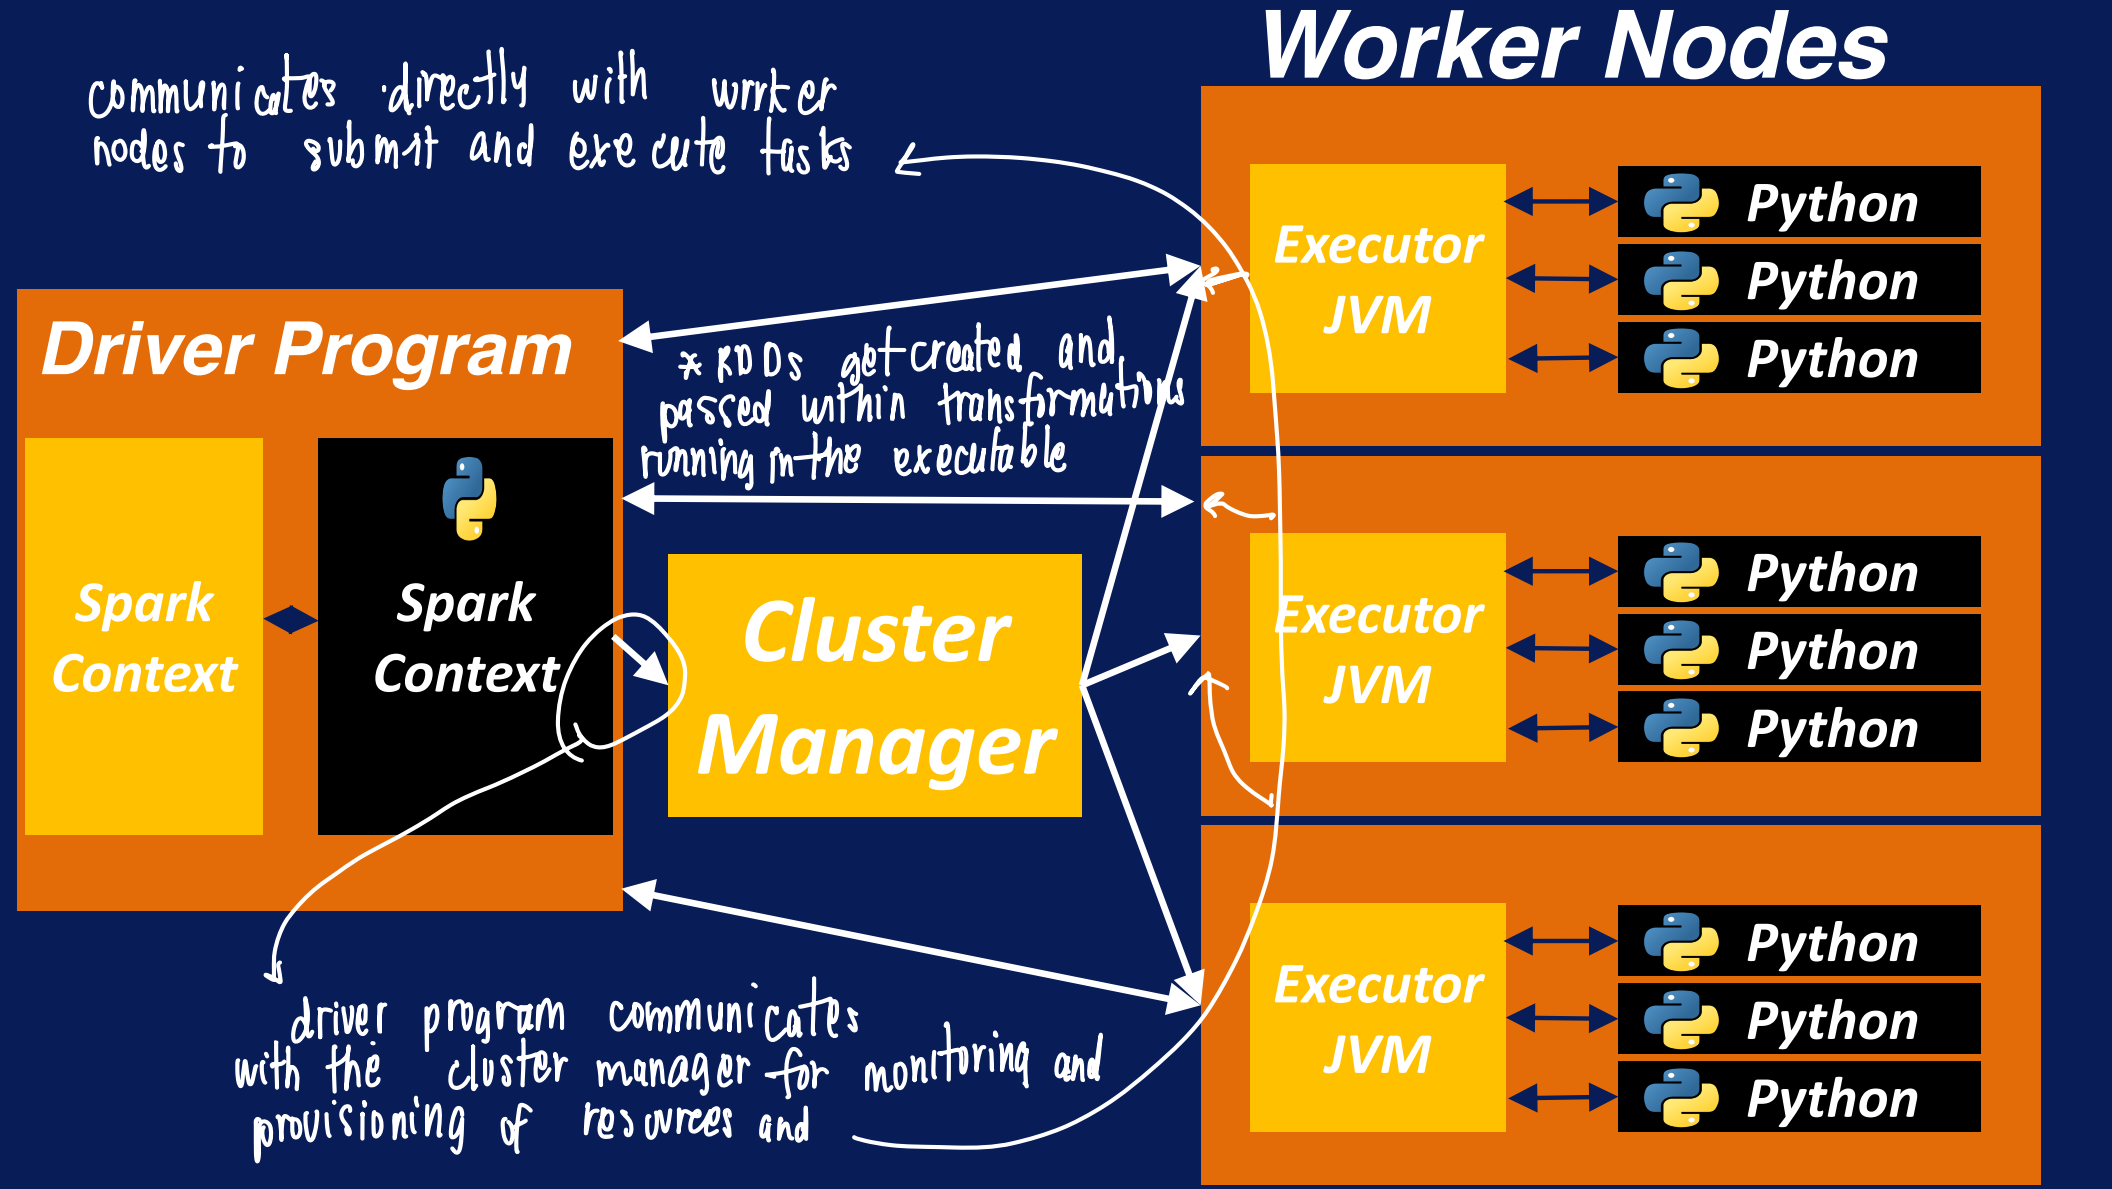
communicates directly with worker nodes to submit and execute tasks
Worker Nodes
Driver Program
Spark Context
Spark Context
Cluster Manager
* RDDs get created and passed within transformations running in the executable
Executor JVM
Python
Python
Python
Executor JVM
Python
Python
Python
Executor JVM
Python
Python
Python
driver program communicates with the cluster manager for monitoring and provisioning of resources and

Cloudera VM

everything runs locally
(a single machine)
- with driver program
- executor JVM
- and single pyspark process