

# Big Data Management: The “M” in DBMS

when to use  
?  $\nwarrow$  hadoop  
parallel

**SDSC** SAN DIEGO  
SUPERCOMPUTER CENTER



# DBMS-based and non-DBMS-based Approaches to Big Data



**SDSC** SAN DIEGO  
SUPERCOMPUTER CENTER

# After this video you will be able to

- Explain the various advantages of using a DBMS over a file system
- Specify the differences between a parallel and distributed file system
- Briefly describe a MapReduce-style DBMS

DBMS vs file system  
parallel vs DBMS

# Storing Data – Files vs. DBMS

- In the old times, database operations were applications in file systems
- Problems
  - Data redundancy, inconsistency and isolation — *duplication or inconsistencies (hard to determine)*
  - Each task a program  $\begin{cases} \text{data access} \\ \text{data update} \end{cases}$
  - Data integrity — *constraints : condition as part of condition of program*
  - Atomicity of updates
    - └ *changes as single unit (altogether) — everything or nothing*

# Advantages of a DBMS

→ state what we want without saying how

- **Declarative query languages**

- No more task-based programs

- **Data independence** : isolate users from record  
or long as logic is clear ( tables & attributes )

- Applications don't worry about data storage formats and locations

- **Efficient access through optimization**

- The system automatically finds an efficient way to access data

powerful ← data structures  
algorithms  
principles

# Advantages of a DBMS

- **Data integrity and security**

- Methods to keep the accuracy and consistency of data despite failure

- ACID properties of transactions
- Failure recovery

- **Concurrent access**

- Many users can simultaneously access data without conflict

• systems fail ! + malicious process  
• safety and failure recovery

(single transactions) — even if it has multiple changes  
properties ACID:

- atomicity
- consistency: data valid (rules, constraint)
- isolation: concurrency (multiple people) update simultaneously
- durability: data remains even after power loss, crashes, errors


serially  
eg  
|  
selling  
last ticket



# Parallel and Distributed DBMS

- Parallel database system

- Improve performance through parallel implementation
- Often allows data replication
  - Data redundancy against table corruption
  - More concurrent queries

eg: 

- tables spread across machines
- operations use parallel algorithms
- allow replications
  - └ data redundancy
  - └ failure of replicas
  - └ replicas synched

- Distributed database system: !

- Data is stored across several sites, each site managed by a DBMS capable of running independently !

*Does your big data problem  
need these facilities?*

answer  
can be  
negative

- network of independently running DBMS that communicate with each other
- one component knows some part of the schema of its neighbor in DBMS and can pass query to

# DBMS and MapReduce-style Systems

- Started with a different problem focus

- DBMSs: efficient storage, transactions and retrieval

- Partitioned data parallelism — different parts of logical table can physically reside on different machines
- Account for computation and communication cost \$\$
- Not node failure — classical <sup>DBMS</sup> did not take into account failures

- Mapreduce-style systems: complex data processing over a cluster of machines

- HDFS-based
- Analytics – data mining, clustering, machine learning
- Multi-stage, problem-specific algorithms — <sup>hard to</sup> develop in standard
- Operate on wider variety of data including text

— originally developed not for storage and retrieval, but for distributive processing of large amounts of data

- number of machines could go up
- issues automatically accounted for (like node failure)
- complex applications, like <sup>data mining</sup> data clustering



# Shifting Requirements

— tension points

- **Data loading – a new bottleneck**

— analysis on the data must be performed within a given time after it's arrival

- Does the application need data sooner than the loading time?

- **Too much functionality**

- Does the application use only a few data management features?

- **Combined Transactional and Analytical Capabilities**

eg. real time decision support

— optimization < support for efficient analytical operations meets transactional guarantees

# No Single Solution

combination of traditional requirements → new capabilities and products in the big data management

## • Mixed solutions

- DBMS on HDFS — new techniques that use MR → side door for MR-style operations
  - Hadoop-DBMS interoperation → flexibility to use both forms of data processing
- Relational operations in MapReduce systems like Spark ← eg
- Streaming input to DBMS — large distributed management operations to accept streaming data → design: analysis known before to perform
  - ↳ as new data records arrive, keep record of data in memory to finish computation
- New parallel programming models for analytical computation within DBMS
  - ↳ large scale distributed algorithms emerge to solve analytics problems
  - ↳ MR style algorithms
  - ↳ evolve